

Task 1: Linear Regression

Goal: Predict house prices based on various features.

- **Dataset:** Use the "California Housing Prices" dataset (or a similar dataset with features like number of bedrooms, bathrooms, square footage, etc.).
- **Instructions:**
 - Load the dataset and explore it using **Pandas**.
 - Visualize the correlation between the features and the target variable (house prices) using **Seaborn heatmaps**.
 - Split the data into training and testing sets.
 - Build a **Linear Regression model** to predict house prices.

Task 2: Logistic Regression

Goal: Classify whether a patient has diabetes based on health parameters.

- **Dataset:** Use the **Pima Indians Diabetes dataset** (available in libraries like sklearn).
- **Instructions:**
 - Load and preprocess the dataset, handling any missing values.
 - Perform **Exploratory Data Analysis (EDA)** to understand feature distributions and correlations.
 - Split the data into training and testing sets.
 - Train a **Logistic Regression model** to classify patients as diabetic or non-diabetic.

Task 3: Decision Tree

Goal: Predict whether a student will pass or fail based on study habits and attendance.

- **Dataset:** Use a dataset containing student information (you can create a simple dataset with features like hours studied, attendance percentage, and past grades).
- **Instructions:**
 - Load and preprocess the dataset.
 - Visualize feature importance using bar plots.
 - Build a **Decision Tree classifier** to predict if a student will pass or fail.
 - Evaluate the model using metrics such as **accuracy**.

Task 4: Random Forest

Goal: Classify whether a customer will buy a product based on their browsing history and demographic information.

- **Dataset:** Use a **Customer Purchase Behavior dataset** (or create your own dataset with features like age, gender, browsing time, product categories visited, etc.).
- **Instructions:**
 - Load and preprocess the data, including feature scaling if necessary.
 - Split the dataset into training and testing sets.
 - Train a **Random Forest classifier** to predict customer purchases.
 - Tune hyperparameters (like the number of trees, max depth) to improve the model's performance.